

# MongoDB

Christof Strauch "NoSQL Databases"

Podkopaev Anton, podkoav239@gmail.com

SE, SPbSU

28 October 2013

# MongoDB

- Schema-free document database
- Written in C++
- Open-source, 10gen Inc.
- Adjective *humongous*
- **Goal:**  
Gap between key-/value-stores and RDBMSs
- **Users:**  
SourceForge.net, foursquare, the New York Times, bit.ly, DIASPORA, Yandex, etc.

# DB structure

- DBs reside on MongoDB server
- Server hosts many of such DBs
- DBs contain collections
- Collections consist of documents

# Collection

- **Def:** named groupings of documents
- A collection may be heterogeneous
- Manual suggests:  
*One database collection for each of your top level objects*

# Creation

- Created automatically at the first document is inserted into a DB
- Parameters
  - auto-indexing
  - preallocated disk space
  - size limits
  - ...

- Explicit creation

```
db.createCollection(  
    "mycoll", {size: 100000000, autoIndexId: false}  
)
```

# Namespaces

- Hierarchical namespaces using a dot-notation  
wiki.articles  
wiki.categories  
wiki.authors
- The collection namespace is flat from the DB's perspective

# Document

- The abstraction and unit of data storable in MongoDB
- Comparable to
  - XML doc
  - JSON doc
  - Python dic
  - Ruby hash
- In fact, the format is called BSON
- 4Mb — limit of document size

# Example

```
{  
  title: "MongoDB",  
  , last_editor: "172.5.123.91"  
  , last_modified: new Date("9/23/2010")  
  , body: "MongoDB is a..."  
  , categories:  
    ["Database"  
    , "NoSQL"  
    , "Document Database"  
    ]  
  , reviewed: false  
}
```



# Datatypes

- boolean, integer, double
- string (UTF-8), regular expressions, code (JS)
- object (BSON)
- object id (12 bytes)
  - timestamp (4 bytes)
  - Machine id (3 bytes)
  - MongoDB process id (2 bytes)
  - counter (3 bytes)
- null
- array
- date

# References

- Doesn't provide a foreign key mechanism
- Manual assignment of ref field the value of the *\_id*
- DBRef
  - Can be dereferenced automatically
  - { \$ref: <collectionName>, \$id: <documentId> }

# References vs Embeddings

- Criteria for Object References
  - First-class domain objects
  - Many-to-many
  - Queried in large numbers
  - Large objects
- Criteria for Object Embeddings
  - "Line-item detail" characteristic
  - Aggregation relationship
  - Not referenced by another object
  - Request performance

# Selection

- Equivalent to the WHERE clause in SQL
- In case of empty object, all documents of a collection are returned

## Example

```
db.collectionName.find( { title: "MongoDB" } )
```

## General form

```
<fieldname>: { $<operator>: <value> }
```

```
<fieldname>: { $<operator>: <value>, $<operator>: <value> }  
// AND-junction
```

# Selection.Operators

```
\$ne // !=
\$gt, \$gte, \$lt, \$lte // >, >=, <, <=
{ \$mod: [2, 1] }
{ \$in : ["NoSQL", "Document Databases"] }
{ \$nin: ["NoSQL", "Document Databases"] }
{ \$all: ["NoSQL", "Document Databases"] }
{ \$size: 2 } // for array
{ \$exists: false }

// OR
{$or: [{ $reviewed: {$exists: true}, {$categories: {$size: 2}}
$nor
$not

{ field.index.field: value }

db.collection.find({ $where: "this.a==1" })
```

# Projection

```
db.collection.find(  
    {<selection criteria>}, {<f1>: 1, <f2>: 1, ...}  
)  
db.collection.find(  
    {<selection criteria>}, {<f1>: 0, <f2>: 0, ...}  
)
```

# Result processing

- sort
- limit
- skip
- count

```
db.collection.find({<selection>},  
                  {inclusion/exclusion},  
                  <limit>, <skip>  
)
```

```
db.collection(...).count()
```

# Cursors

```
var cursor = db.collection.find(...)  
cursor.forEach( function(result) {...} )  
  
db.collection.findOne( {_id: 239239} )
```



## Query optimizer

- Supports ad-hoc queries
- Not based on statistics
- Not model the costs of multiple possible query plans
- Execute different query plans in parallel
- Stops all of them as soon as the first has finished
- The system is non-relational
- There are no joins
- The space of possible query plans much smaller